

# models

December 7, 2023

```
[33]: import pickle
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from nltk.tokenize import RegexpTokenizer
from gensim.models import Word2Vec
from nltk.tokenize import word_tokenize
from tensorflow import keras
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from tensorflow.keras.layers import Input, Dense
```

## 1 Generic Fucntions

```
[2]: labels=pickle.load( open('labels.pkl', 'rb'))
features=pickle.load( open('features.pkl', 'rb'))
features_tfidf= pickle.load( open('features_tfidf.pkl', 'rb'))
word_tfidf_weights=pickle.load( open('word_tfidf_weights.pkl', 'rb'))
non_text_features_np = features.drop(columns=['title', 'text', 'combined_text', '
↳ 'label']).to_numpy()
```

```
[3]: def get_metrics(clf,test_ft,test_labels):
    pred_labels=clf.predict(test_ft)
    ↳
    ↳print(classification_report(pred_labels,test_labels,target_names=["fake","real"]))
```

```
[4]: def plot_ROC(test):

    pred_labels=clf.predict(test_ft)
    fpr, tpr, thresholds = roc_curve(test,pred_labels)
    roc_auc = auc(fpr, tpr)
```

```

# Plot ROC curve
plt.plot(fpr, tpr, label='ROC curve (AUC = {:.2f})'.format(roc_auc))
plt.plot([0, 1], [0, 1], 'k--') # Random guessing line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc='lower right')
plt.show()

```

```

[101]: class ANN_model(tf.keras.models.Model):
        def __init__(self):
            super(ANN_model, self).__init__()
            self.layer1= Dense(64, activation="relu")
            self.layer2= Dense(32, activation="relu")
            self.label= Dense(1, activation="sigmoid")

        def call(self, inputs):
            x=self.layer1(inputs)
            x=self.layer2(x)
            x=self.label(x)

            return x

```

```

[8]: def plot_loss_acc(history):

    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()

    plt.plot(history.history['accuracy'], label='Training Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.show()

```

```

[10]: def split_data_ml(features, labels):

    scaler=StandardScaler()
    norm_features=scaler.fit_transform(features_tfidf)

```

```

    ↪
    ↪train_ft,test_ft,train_labels,test_labels=train_test_split(norm_features,labels,test_size=0
    ↪2, train_size=0.8)
    print("-----Train data-----\n")
    print(f" train data features shape {train_ft.shape} \n train data labels_
    ↪shape {train_labels.shape}\n")
    print("-----Test data-----\n")
    print(f" test data features shape {test_ft.shape} \n test data labels shape_
    ↪{test_labels.shape}")

    return train_ft,test_ft,train_labels,test_labels

```

```

[30]: def split_data_dl(features,labels):

    ↪
    ↪train_ft,test_ft,train_labels,test_labels=train_test_split(norm_features,labels,test_size=0
    ↪2, train_size=0.8)
    X_train, X_val, y_train, y_val = train_test_split(train_ft, train_labels,
    ↪test_size=0.2, random_state=42)

    print("Training set shape:", X_train.shape)
    print("Validation set shape:", X_val.shape)
    print("Test set shape:", test_ft.shape)

    return X_train,y_train,X_val,y_val,test_ft,test_labels

```

## 2 1. TF-IDF vectors + text characteristics

### 2.1 Machine Learning models

```

[21]: features_tfidf_txt = np.hstack((features_tfidf, non_text_features_np))

```

```

[46]: pickle.dump(features_tfidf_txt, open('features_tfidf_txt.pkl', 'wb'))
    features_tfidf_txt = pickle.load( open('features_tfidf_txt.pkl', 'rb'))

```

#### 2.1.1 Logistic Regression

```

[22]: train_ft,test_ft,train_labels,test_labels=split_data_ml(features_tfidf_txt,labels)

```

```

-----Train data-----

train data features shape (57229, 115)
train data labels shape (57229,)

-----Test data-----

```


```
test data features shape (14308, 115)
test data labels shape (14308,)
```

```
[23]: log_clf = LogisticRegression(random_state=0).fit(train_ft, train_labels)
```

```
[24]: get_metrics(log_clf, test_ft, test_labels)
```

	precision	recall	f1-score	support
fake	0.90	0.93	0.91	6762
real	0.93	0.91	0.92	7546
accuracy			0.92	14308
macro avg	0.92	0.92	0.92	14308
weighted avg	0.92	0.92	0.92	14308

```
[25]: plot_ROC(test_labels)
```




### 2.1.2 Linear Discriminant Analysis

```
[26]: clf = LinearDiscriminantAnalysis().fit(train_ft, train_labels)
```

```
[27]: get_metrics(clf, test_ft, test_labels)
```

	precision	recall	f1-score	support
fake	0.88	0.93	0.90	6580
real	0.94	0.89	0.91	7728
accuracy			0.91	14308
macro avg	0.91	0.91	0.91	14308
weighted avg	0.91	0.91	0.91	14308

```
[28]: plot_ROC(test_labels)
```



### 2.1.3 Artificial Neural network

```
[31]: x_train,y_train,X_val,y_val,test_ft,test_labels=split_data_dl(features_tfidf_txt,labels)
```

```
Training set shape: (45783, 115)
Validation set shape: (11446, 115)
Test set shape: (14308, 115)
```

```
[34]: tfidf_txt_model= ANN_model()
      tfidf_txt_model.build((None,115))
      print(tfidf_txt_model.summary())
```

```
Model: "ann_model_1"
```

Layer (type)	Output Shape	Param #
dense (Dense)	multiple	7424
dense_1 (Dense)	multiple	2080
dense_2 (Dense)	multiple	33

```
=====  
Total params: 9,537  
Trainable params: 9,537  
Non-trainable params: 0
```

```
-----  
None
```

```
[102]: tb_callback=tf.keras.callbacks.TensorBoard(log_dir="log/",histogram_freq=1)
      tfidf_txt_model.compile(optimizer='adam', loss='binary_crossentropy',  
        ↪metrics=['accuracy'])  
      history=tfidf_txt_model.fit(train_ft, train_labels, epochs=10, batch_size=32,  
        ↪validation_data=(X_val, y_val),callbacks=[tb_callback])
```

```
Epoch 1/10
```

```
1789/1789 [=====] - 4s 2ms/step - loss: 0.0969 -  
accuracy: 0.9642 - val_loss: 0.0852 - val_accuracy: 0.9678
```

```
Epoch 2/10
```

```
1789/1789 [=====] - 3s 2ms/step - loss: 0.0854 -  
accuracy: 0.9678 - val_loss: 0.0751 - val_accuracy: 0.9727
```

```
Epoch 3/10
```

```
1789/1789 [=====] - 3s 2ms/step - loss: 0.0783 -  
accuracy: 0.9713 - val_loss: 0.0723 - val_accuracy: 0.9754
```

```
Epoch 4/10
```

```
1789/1789 [=====] - 3s 2ms/step - loss: 0.0729 -  
accuracy: 0.9736 - val_loss: 0.0692 - val_accuracy: 0.9746
```

```
Epoch 5/10
```

```

1789/1789 [=====] - 3s 2ms/step - loss: 0.0680 -
accuracy: 0.9752 - val_loss: 0.0716 - val_accuracy: 0.9765
Epoch 6/10
1789/1789 [=====] - 3s 2ms/step - loss: 0.0627 -
accuracy: 0.9765 - val_loss: 0.0756 - val_accuracy: 0.9716
Epoch 7/10
1789/1789 [=====] - 3s 2ms/step - loss: 0.0581 -
accuracy: 0.9789 - val_loss: 0.0670 - val_accuracy: 0.9779
Epoch 8/10
1789/1789 [=====] - 3s 2ms/step - loss: 0.0550 -
accuracy: 0.9797 - val_loss: 0.0693 - val_accuracy: 0.9766
Epoch 9/10
1789/1789 [=====] - 3s 2ms/step - loss: 0.0502 -
accuracy: 0.9817 - val_loss: 0.0636 - val_accuracy: 0.9803
Epoch 10/10
1789/1789 [=====] - 4s 2ms/step - loss: 0.0476 -
accuracy: 0.9826 - val_loss: 0.0608 - val_accuracy: 0.9812

```

```

[38]: # Evaluate the model on the test set
test_loss, test_accuracy = tfidf_txt_model.evaluate(test_ft, test_labels)
print(f'Test Loss: {test_loss}, Test Accuracy: {test_accuracy}')

```

```


448/448 [=====] - 1s 2ms/step - loss: 0.0924 -
accuracy: 0.9683
Test Loss: 0.09243853390216827, Test Accuracy: 0.9683393836021423

```


```

[39]: plot_loss_acc(history)

```







```
[100]: tfidf_txt_model.save_weights("tfidf_txt_model.h5")
```

### 3 2. Word2vec vectors + text characteristics

#### 4 Machine Learning Models

```
[40]: doc=list(features["combined_text"])
```

```
[41]: def tokenize_words(text):
    tokenized_words,vocab=[], []
    for i in text:
        tokenizer = RegexpTokenizer(r'\w+')          #using regular expression
        wrd_token=tokenizer.tokenize(i)
        tokenized_words.append(wrd_token)
        vocab.extend(wrd_token)
    return tokenized_words, vocab

doc_tokens_w2v,vocab=tokenize_words(doc)
```

```
[42]: model_doc=Word2Vec(doc_tokens_w2v,min_count=2,window=4)
```

```
[43]: sentence_vectors = []
for sentence in doc:
    tokens = word_tokenize(sentence.lower())
    sentence_vector = np.zeros(model_doc.vector_size)
    if len(tokens)==0:
        sentence_vectors.append([0]*100)
    else:
        for token in tokens:
            if token in model_doc.wv:
                token_vector = model_doc.wv[token]
                sentence_vector += token_vector
        sentence_vector /= len(tokens)
        sentence_vectors.append(sentence_vector)

# Convert the list of sentence vectors to a NumPy array
sentence_vectors = np.array(sentence_vectors)
```

```
[44]: features_w2v_txt = np.hstack((sentence_vectors, non_text_features_np))
```

```
[47]: pickle.dump(features_w2v, open('features_w2v_txt.pkl', 'wb'))
```

```
[48]: features_w2v_txt = pickle.load( open('features_w2v_txt.pkl', 'rb'))
```

```
[49]: train_ft,test_ft,train_labels,test_labels=split_data_ml(features_w2v_txt,labels)
```

-----Train data-----

```
train data features shape (57229, 115)
train data labels shape (57229,)
```

-----Test data-----


```
test data features shape (14308, 115)
test data labels shape (14308,)
```

#### 4.0.1 Logistic Regression

```
[50]: clf = LogisticRegression(random_state=0).fit(train_ft, train_labels)
get_metrics(clf,test_ft,test_labels)
plot_ROC(test_labels)
```

	precision	recall	f1-score	support
fake	0.91	0.92	0.91	6897
real	0.93	0.91	0.92	7411


accuracy			0.92	14308
macro avg	0.92	0.92	0.92	14308
weighted avg	0.92	0.92	0.92	14308



```
[51]: get_metrics(clf,test_ft,test_labels)
```

	precision	recall	f1-score	support
fake	0.91	0.92	0.91	6897
real	0.93	0.91	0.92	7411
accuracy			0.92	14308
macro avg	0.92	0.92	0.92	14308
weighted avg	0.92	0.92	0.92	14308

```
[52]: plot_ROC(test_labels)
```




#### 4.0.2 Linear Discriminant Analysis

```
[53]: clf = LinearDiscriminantAnalysis().fit(train_ft, train_labels)
```

```
[54]: get_metrics(clf, test_ft, test_labels)
```

	precision	recall	f1-score	support
fake	0.88	0.92	0.90	6741
real	0.93	0.89	0.91	7567
accuracy			0.91	14308
macro avg	0.90	0.91	0.91	14308
weighted avg	0.91	0.91	0.91	14308

```
[55]: plot_ROC(test_labels)
```



### 4.0.3 Artificial Neural Network

```
[56]: x_train,y_train,X_val,y_val,test_ft,test_labels=split_data_dl(features_w2v_txt,labels)
```

```
Training set shape: (45783, 115)
Validation set shape: (11446, 115)
Test set shape: (14308, 115)
```

```
[57]: tfidf_w2v_model= ANN_model()
tfidf_w2v_model.build((None,115))
print(tfidf_w2v_model.summary())
```

```
Model: "ann_model_2"
```

Layer (type)	Output Shape	Param #
dense_3 (Dense)	multiple	7424
dense_4 (Dense)	multiple	2080

dense\_5 (Dense)                      multiple                      33

```
=====
Total params: 9,537
Trainable params: 9,537
Non-trainable params: 0
```

```
-----
None
```

```
[58]: tfidf_w2v_model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪ metrics=['accuracy'])
history=tfidf_w2v_model.fit(train_ft, train_labels, epochs=10, batch_size=32,
    ↪ validation_data=(X_val, y_val))
```


```
Epoch 1/10
1789/1789 [=====] - 6s 3ms/step - loss: 0.2207 -
accuracy: 0.9081 - val_loss: 0.1626 - val_accuracy: 0.9355
Epoch 2/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1591 -
accuracy: 0.9368 - val_loss: 0.1464 - val_accuracy: 0.9393
Epoch 3/10
1789/1789 [=====] - 6s 3ms/step - loss: 0.1415 -
accuracy: 0.9429 - val_loss: 0.1313 - val_accuracy: 0.9472
Epoch 4/10
1789/1789 [=====] - 6s 3ms/step - loss: 0.1284 -
accuracy: 0.9493 - val_loss: 0.1191 - val_accuracy: 0.9540
Epoch 5/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1188 -
accuracy: 0.9535 - val_loss: 0.1127 - val_accuracy: 0.9553
Epoch 6/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1096 -
accuracy: 0.9575 - val_loss: 0.1087 - val_accuracy: 0.9597
Epoch 7/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1011 -
accuracy: 0.9612 - val_loss: 0.1018 - val_accuracy: 0.9640
Epoch 8/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.0949 -
accuracy: 0.9636 - val_loss: 0.0984 - val_accuracy: 0.9651
Epoch 9/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.0885 -
accuracy: 0.9661 - val_loss: 0.0922 - val_accuracy: 0.9676
Epoch 10/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.0816 -
accuracy: 0.9700 - val_loss: 0.0900 - val_accuracy: 0.9702
```


```
[60]: # Evaluate the model on the test set
test_loss, test_accuracy = tfidf_w2v_model.evaluate(test_ft, test_labels)
```

```
print(f'Test Loss: {test_loss}, Test Accuracy: {test_accuracy}')
```

```
448/448 [=====] - 1s 2ms/step - loss: 0.0943 -  
accuracy: 0.9678  
Test Loss: 0.09433166682720184, Test Accuracy: 0.9677802920341492
```

```
[61]: plot_loss_acc(history)
```





```
[63]: tfidf_w2v_model.save_weights("tfidf_w2v_model.h5")
```

## 5 3. TF-IDF weighted Word2vec vectors

### 5.1 Machine learning models

```
[64]: sentence_vectors_tf = []

# Process each sentence and compute its vector
for i,sentence in enumerate(doc):
    tokens = word_tokenize(sentence.lower())
    sentence_vector = np.zeros(model_doc.vector_size)
    if len(tokens)==0:
        sentence_vectors_tf.append([0]*100)
    else:
        for token in tokens:
            if token in model_doc.wv:
                token_vector = model_doc.wv[token]
                sentence_vector += token_vector
```



```

        sentence_vector /= len(tokens)
        sentence_vector *= word_tfidf_weights[i]
        sentence_vectors_tf.append(sentence_vector)

# Convert the list of sentence vectors to a NumPy array
sentence_vectors_tf = np.array(sentence_vectors)
sentence_vectors_tf = np.nan_to_num(sentence_vectors_tf, nan=0.0)

```

```
[73]: sentence_vectors_tf.shape
```

```
[73]: (71537, 100)
```

```
[77]: pickle.dump(sentence_vectors_tf, open('features_w2v_tf_wts.pkl', 'wb'))
```

```
[78]: features_w2v_tf_wts=pickle.load( open('features_w2v_tf_wts.pkl', 'rb'))
```

```
[80]: train_ft,test_ft,train_labels,test_labels=split_data_ml(features_w2v_tf_wts,labels)
```

```
-----Train data-----
```

```

train data features shape (57229, 115)
train data labels shape (57229,)

```

```
-----Test data-----
```

```

test data features shape (14308, 115)
test data labels shape (14308,)

```


### 5.1.1 Logistic Regression

```

[81]: clf = LogisticRegression(random_state=0, max_iter=300).fit(train_ft,
    ↪train_labels)
get_metrics(clf,test_ft,test_labels)
plot_ROC(test_labels)

```


	precision	recall	f1-score	support
fake	0.90	0.92	0.91	6816
real	0.93	0.91	0.92	7492
accuracy			0.92	14308
macro avg	0.92	0.92	0.92	14308
weighted avg	0.92	0.92	0.92	14308



### 5.1.2 Linear Discriminant Analysis

```
[82]: clf = LinearDiscriminantAnalysis().fit(train_ft, train_labels)
      get_metrics(clf, test_ft, test_labels)
      plot_ROC(test_labels)
```

	precision	recall	f1-score	support
fake	0.88	0.92	0.90	6658
real	0.93	0.89	0.91	7650
accuracy			0.91	14308
macro avg	0.91	0.91	0.91	14308
weighted avg	0.91	0.91	0.91	14308



### 5.1.3 Artificial Neural Networks

```
[83]: x_train,y_train,X_val,y_val,test_ft,test_labels=split_data_dl(features_w2v_tf_wts,labels)
```

Training set shape: (45783, 115)

Validation set shape: (11446, 115)

Test set shape: (14308, 115)

```
[84]: tfidf_w2v_wts_model= ANN_model()
tfidf_w2v_wts_model.build((None,115))
print(tfidf_w2v_model.summary())
```

Model: "ann\_model\_3"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	multiple	7424
dense_7 (Dense)	multiple	2080

dense\_8 (Dense)

multiple

33

```
=====
Total params: 9,537
Trainable params: 9,537
Non-trainable params: 0
-----
```

None

```
[85]: opt = keras.optimizers.SGD(learning_rate=0.1)
tfidf_w2v_wts_model.compile(optimizer=opt, loss='binary_crossentropy',
    ↪metrics=['accuracy'])
history=tfidf_w2v_wts_model.fit(train_ft, train_labels, epochs=10,
    ↪batch_size=32, validation_data=(X_val, y_val))
```

Epoch 1/10

1789/1789 [=====] - 7s 3ms/step - loss: 0.2162 - accuracy: 0.9123 - val\_loss: 0.1794 - val\_accuracy: 0.9275

Epoch 2/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1643 - accuracy: 0.9348 - val\_loss: 0.1585 - val\_accuracy: 0.9368

Epoch 3/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1486 - accuracy: 0.9409 - val\_loss: 0.1460 - val\_accuracy: 0.9416

Epoch 4/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1369 - accuracy: 0.9464 - val\_loss: 0.1368 - val\_accuracy: 0.9464

Epoch 5/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1271 - accuracy: 0.9498 - val\_loss: 0.1314 - val\_accuracy: 0.9469

Epoch 6/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1191 - accuracy: 0.9533 - val\_loss: 0.1231 - val\_accuracy: 0.9519

Epoch 7/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1123 - accuracy: 0.9565 - val\_loss: 0.1176 - val\_accuracy: 0.9552

Epoch 8/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1050 - accuracy: 0.9592 - val\_loss: 0.1167 - val\_accuracy: 0.9553

Epoch 9/10

1789/1789 [=====] - 5s 3ms/step - loss: 0.1015 - accuracy: 0.9607 - val\_loss: 0.1078 - val\_accuracy: 0.9594


Epoch 10/10


1789/1789 [=====] - 5s 3ms/step - loss: 0.0955 - accuracy: 0.9633 - val\_loss: 0.1046 - val\_accuracy: 0.9615

```
[86]: # Evaluate the model on the test set
test_loss, test_accuracy = tfidf_w2v_wts_model.evaluate(test_ft, test_labels)
print(f'Test Loss: {test_loss}, Test Accuracy: {test_accuracy}')
```

```
448/448 [=====] - 1s 2ms/step - loss: 0.0964 -
accuracy: 0.9639
Test Loss: 0.09637467563152313, Test Accuracy: 0.9638663530349731
```

```
[87]: plot_loss_acc(history)
```





```
[88]: tfidf_w2v_wts_model.save_weights("tfidf_w2v_wts_model.h5")
```

## 6 4 TF-IDF weighted Word2vec with non numeric features

### 6.1 Machine Learning Models

```
[89]: features_w2v_tfidf_txt = np.hstack((sentence_vectors_tf, non_text_features_np))
```

#### 6.1.1 Logistic Regression

```
[90]: train_ft, test_ft, train_labels, test_labels = split_data_ml(features_w2v_tfidf_txt, labels)
```

```
-----Train data-----
```

```
train data features shape (57229, 115)
```

```
train data labels shape (57229,)
```


```
-----Test data-----
```

```
test data features shape (14308, 115)
```

test data labels shape (14308,)

```
[91]: clf = LogisticRegression(random_state=0, max_iter=300).fit(train_ft, train_labels)
      get_metrics(clf, test_ft, test_labels)
      plot_ROC(test_labels)
```


	precision	recall	f1-score	support
fake	0.90	0.92	0.91	6809
real	0.93	0.91	0.92	7499
accuracy			0.91	14308
macro avg	0.91	0.92	0.91	14308
weighted avg	0.92	0.91	0.91	14308



### 6.1.2 Linear Discriminant Analysis

```
[92]: clf = LinearDiscriminantAnalysis().fit(train_ft, train_labels)
      get_metrics(clf, test_ft, test_labels)
      plot_ROC(test_labels)
```

	precision	recall	f1-score	support
fake	0.88	0.93	0.90	6666
real	0.93	0.89	0.91	7642
accuracy			0.91	14308
macro avg	0.91	0.91	0.91	14308
weighted avg	0.91	0.91	0.91	14308



### 6.1.3 Artificial Neural Networks

```
[93]: x_train, y_train, X_val, y_val, test_ft, test_labels = split_data_dl(features_w2v_tfidf_txt, labels)
```

Training set shape: (45783, 115)



Validation set shape: (11446, 115)  
Test set shape: (14308, 115)

```
[95]: tfidf_w2v_txt_model= ANN_model()  
      tfidf_w2v_txt_model.build((None,115))  
      print(tfidf_w2v_txt_model.summary())
```

Model: "ann\_model\_6"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	multiple	7424
dense_16 (Dense)	multiple	2080
dense_17 (Dense)	multiple	33

=====  
Total params: 9,537  
Trainable params: 9,537  
Non-trainable params: 0  
=====  
None

```
[96]: opt = keras.optimizers.SGD(learning_rate=0.1)  
      tfidf_w2v_txt_model.compile(optimizer=opt, loss='binary_crossentropy',  
      ↪metrics=['accuracy'])  
      history=tfidf_w2v_txt_model.fit(train_ft, train_labels, epochs=10,  
      ↪batch_size=32, validation_data=(X_val, y_val))
```

Epoch 1/10  
1789/1789 [=====] - 6s 3ms/step - loss: 0.2148 -  
accuracy: 0.9143 - val\_loss: 0.1763 - val\_accuracy: 0.9291  
Epoch 2/10  
1789/1789 [=====] - 5s 3ms/step - loss: 0.1650 -  
accuracy: 0.9348 - val\_loss: 0.1527 - val\_accuracy: 0.9404  
Epoch 3/10  
1789/1789 [=====] - 5s 3ms/step - loss: 0.1486 -  
accuracy: 0.9411 - val\_loss: 0.1400 - val\_accuracy: 0.9446  
Epoch 4/10  
1789/1789 [=====] - 5s 3ms/step - loss: 0.1369 -  
accuracy: 0.9463 - val\_loss: 0.1325 - val\_accuracy: 0.9483  
Epoch 5/10  
1789/1789 [=====] - 5s 3ms/step - loss: 0.1277 -  
accuracy: 0.9503 - val\_loss: 0.1233 - val\_accuracy: 0.9521  
Epoch 6/10  
1789/1789 [=====] - 5s 3ms/step - loss: 0.1203 -  
accuracy: 0.9530 - val\_loss: 0.1231 - val\_accuracy: 0.9518

```

Epoch 7/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1132 -
accuracy: 0.9559 - val_loss: 0.1113 - val_accuracy: 0.9584
Epoch 8/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1077 -
accuracy: 0.9586 - val_loss: 0.1178 - val_accuracy: 0.9537
Epoch 9/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.1013 -
accuracy: 0.9613 - val_loss: 0.1130 - val_accuracy: 0.9591
Epoch 10/10
1789/1789 [=====] - 5s 3ms/step - loss: 0.0962 -
accuracy: 0.9628 - val_loss: 0.1030 - val_accuracy: 0.9618

```

```

[97]: # Evaluate the model on the test set
test_loss, test_accuracy = tfidf_w2v_txt_model.evaluate(test_ft, test_labels)
print(f'Test Loss: {test_loss}, Test Accuracy: {test_accuracy}')

```

```


448/448 [=====] - 1s 2ms/step - loss: 0.1041 -
accuracy: 0.9605
Test Loss: 0.10413851588964462, Test Accuracy: 0.9605116248130798


```

```

[98]: plot_loss_acc(history)

```





```
[99]: tfidf_w2v_txt_model.save_weights("tfidf_w2v_txt_model.h5")
```

```
[118]: %load_ext tensorboard
      %tensorboard --logdir=log/train
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

<IPython.core.display.HTML object>

```
[ ]:
```